THE REVENUE SYSTEM

AGENT 4

Marketing-Sourced Pipeline

A technical implementation guide for building an AI agent that connects every marketing touch to pipeline revenue — so you know exactly what's working, in dollars, not leads.

A step-by-step technical implementation guide.

Note: the tools used in this implementation guide are used for example purposes.

If you don't use Fibbler or a specific CRM, this simply swap out your tools and replace the API url's accordingly.

What This Agent Does

Most marketing teams report on leads, impressions, and MQLs. None of those numbers tell you what actually drove revenue. The Marketing-Sourced Pipeline agent closes that gap by stitching together every channel touch - from first LinkedIn click to closed deal - into a single contact journey, then using Claude to score which signals actually predict pipeline.

The result is a weekly Slack brief that tells you: how much pipeline was created this week, which channels sourced it, which content showed up in winning journeys, and where you're wasting spend.

Workflow at a Glance

CRM CRM → Contact + Deal data (source, stage, close status)

Fibbler → LinkedIn ad engagement mapped to named accounts

Riverside → Webinar attendance + replay engagement

n8n → Journey stitching + signal reconciliation

Claude → Weighted signal scoring + weekly narrative brief

Slack → Weekly output: pipeline → channel → content breakdown

CRM → Auto-tagged deal records with source attribution

What You Need Before Starting

This agent sits at the intersection of your ad stack, event platform, CRM, and automation layer. Before building, confirm access to each of the following tools and data points.

Tool	Purpose	Required Access
CRM	Contact timeline, deal source, close data	CRM API key (private app, read+write)
Fibbler (or other signal software)	LinkedIn ad-to-account matching	Workspace connected, Fibbler API key
Riverside	Webinar registrant + attendee lists	Riverside API key or CSV export webhook
n8n	Automation orchestration	Self-hosted or cloud, CRM + HTTP nodes
Claude API	Signal scoring + brief generation	Anthropic API key (claude-sonnet-4)
Slack	Weekly output delivery	Slack Bot token, target channel ID

CRM Data Requirements

For attribution to work, CRM contacts must have source tracking enabled. Confirm the following properties exist and are being populated on every new contact:

- **Original Source** — set automatically by CRM (Organic Search, Paid Social, etc.)
- **Original Source Drill-Down 1 & 2** — channel-level detail (e.g. LinkedIn Ads, Google)
- **Latest Source** — most recent touch before conversion
- **Associated Deal IDs** — required to link contact journey to pipeline
- **Lifecycle Stage** — Lead, MQL, SQL, Opportunity, Customer

> **Pro Tip: UTM Hygiene**
> Fibbler and Riverside attribution only work if your landing pages pass UTM parameters into CRM on form submit. Audit your forms before building. Every untracked source creates a gap in the journey map.

Step 1: Pull the Contact Journey from CRM

The foundation of this agent is the contact journey: a chronological log of every tracked interaction a contact had before becoming an opportunity or customer. n8n pulls this from CRM's engagement and contact APIs on a weekly schedule.

1A. Set Up the n8n Trigger

Use a Schedule Trigger node set to run every Monday at 6:00 AM. This gives Claude fresh data before your team's weekly review.

```
// n8n Schedule Trigger
Trigger type: Cron
Expression: 0 6 * * 1   // Every Monday at 6:00 AM
```

1B. Fetch Deals Created in the Last 7 Days

Use an HTTP Request node to call the CRM Deals API. Filter for deals created in the rolling 7-day window with a close date in the future or a Closed Won status.

```
POST https://api.[YOUR-CRM-API-URL]/.com/crm/v3/objects/deals/search
Authorization: Bearer {{$credentials.CRMApiKey}}
Content-Type: application/json

{
  "filterGroups": [{
    "filters": [{
      "propertyName": "createdate",
      "operator": "GTE",
      "value": "{{$now.minus({days: 7}).toMillis()}}"
    }]
  }],
  "properties": ["dealname","amount","pipeline","dealstage","closedate",
                 "CRM_owner_id","hs_analytics_source","associations"],
  "limit": 100
}
```

1C. Fetch Associated Contact Timelines

For each deal returned, use a Loop Over Items node to fetch the associated contacts. Then pull each contact's full engagement timeline using the Engagements API. This returns every email, call, meeting, form submission, and page view tied to that contact.

```
GET https://api.[YOUR-CRM-API-
URL]/.com/engagements/v1/engagements/associated/CONTACT/{{contactId}}/paged
Authorization: Bearer {{$credentials.CRMApiKey}}
```

```
// Returns: engagement type, timestamp, metadata
// Types: EMAIL, CALL, MEETING, NOTE, TASK, FORM_SUBMISSION
```

Watch Out: Rate Limits

CRM's Engagements API is rate-limited at 100 requests per 10 seconds on the free tier. If you have more than 100 deals, add a Wait node (1 second) inside your loop. The paid CRM API tier increases this to 150 req/sec.

Step 2: Pull LinkedIn Engagement from Fibbler

CRM's native LinkedIn attribution only captures clicks on Lead Gen Forms. Fibbler goes further: it matches LinkedIn ad impressions and engagements to named accounts in your CRM, even when the contact never clicked a link. This fills the dark funnel.

2A. Export Fibbler Account Engagement

Fibbler provides a REST API and a CSV export. For automation, use the API endpoint to pull account-level engagement for the same 7-day window as your CRM pull.

```
GET https://api.fibbler.co/v1/account-engagement
Authorization: Bearer {{$credentials.fibblerApiKey}}
?from={{$now.minus({days: 7}).toISODate()}}
&to={{$now.toISODate()}}

// Returns: account domain, ad campaign, impressions, clicks, engagement date
```

2B. Match Fibbler Accounts to CRM Contacts

Use the account domain returned by Fibbler to look up matching companies and associated contacts in CRM. This creates the link between "this account saw our LinkedIn campaign" and "this contact became a deal.

```
GET https://[YOUR-CRM-API-URL]/crm/v3/objects/companies/search
Authorization: Bearer {{$credentials.CRMApiKey}}

{
  "filterGroups": [{
    "filters": [{
      "propertyName": "domain",
      "operator": "EQ",
      "value": "{{fibblerAccountDomain}}"
    }]
  }],
  "properties": ["name","domain","associated_contacts"]
}
```

Why This Matters

LinkedIn is heavily underreported in most CRMs. A contact might see your ad 12 times, never click, then respond to an outbound email 3 weeks later. Without Fibbler, that deal gets attributed to outbound. With Fibbler, you see the full warm-up sequence and weight the channel correctly.

Step 3: Pull Webinar Data from Riverside

Webinar attendance is one of the highest-intent signals in a B2B buying journey. Riverside captures registrant and attendee data including how long they watched, whether they asked questions, and whether they rewatched the replay. This data plugs directly into your contact journey.

3A. Pull Registrant and Attendee Lists

Use Riverside's API or post-event webhook to pull attendance data after each webinar. Map each attendee's email back to a CRM contact.

```
GET https://api.riverside.fm/v1/sessions/{{sessionId}}/attendees
Authorization: Bearer {{$credentials.riversideApiKey}}

// Returns: email, registration_date, attended (bool),
//          watch_duration_seconds, replay_views
```

3B. Write Webinar Engagement to CRM Timeline

For each attendee matched to a CRM contact, write a custom timeline event. This makes webinar attendance visible in the contact's CRM record without requiring a native integration.

```
POST https://api.CRM.com/crm/v3/timeline/events
Authorization: Bearer {{$credentials.CRMApiKey}}

{
  "eventTemplateId": "{{your_webinar_template_id}}",
  "objectId": "{{CRMContactId}}",
  "tokens": {
    "webinar_title": "{{sessionTitle}}",
    "attended": "{{attended}}",
    "watch_duration": "{{watchDurationSeconds}}",
    "replay_views": "{{replayViews}}"
  }
}
```

Setup Required: Create a Timeline Event Template

Before writing webinar events, create a custom Timeline Event Template in CRM under Settings > Objects > Contacts > Timeline Events. Define tokens for webinar_title, attended, watch_duration, and replay_views. Copy the eventTemplateId into your n8n workflow.

Step 4: Stitch the Contact Journey in n8n

With CRM engagements, Fibbler ad signals, and Riverside webinar data all pulled, n8n now merges them into a single chronological journey object per deal. This is the data payload that Claude will score.

4A. Merge and Sort by Timestamp

Use a Function node to combine engagement arrays from all three sources, deduplicate by contact ID + timestamp, and sort chronologically.

```
// n8n Function node: Stitch journey
const journeys = {};

for (const engagement of CRMEngagements) {
  const id = engagement.contactId;
  if (!journeys[id]) journeys[id] = [];
  journeys[id].push({
    source: 'CRM',
    type: engagement.type,
    timestamp: engagement.createdAt,
    meta: engagement.metadata
  });
}

for (const signal of fibblerSignals) {
  const id = signal.matchedContactId;
  if (id && journeys[id]) {
    journeys[id].push({
      source: 'linkedin_fibbler',
      type: 'ad_impression',
      timestamp: signal.engagementDate,
      meta: { campaign: signal.campaign, impressions: signal.impressions }
    });
  }
}

// Sort each journey chronologically
for (const id of Object.keys(journeys)) {
  journeys[id].sort((a, b) => new Date(a.timestamp) - new Date(b.timestamp));
}

return [{ json: { journeys, deals: $('CRM Deals').all() } }];
```

Step 5: Claude Scores Channels and Builds the Brief

With stitched journey data in hand, n8n sends a structured prompt to Claude. Claude's job is to do two things: (1) build a weighted signal model that identifies which channels and content pieces appear most in pipeline-generating journeys, and (2) write a human-readable weekly brief.

5A. Construct the Claude Prompt

Pass the full journeys and deal outcomes as JSON context. The prompt instructs Claude to score channels by pipeline contribution, not lead volume.

```
POST https://api.anthropic.com/v1/messages
x-api-key: {{$credentials.claudeApiKey}}
anthropic-version: 2023-06-01
Content-Type: application/json

{
  "model": "claude-sonnet-4-20250514",
  "max_tokens": 2000,
  "messages": [{
    "role": "user",
    "content": "You are a B2B revenue analyst. Analyze these contact journeys and
deal outcomes.
    For each closed-won deal, trace which channels and content appeared in the
journey.
    Build a channel scoring model weighted by pipeline dollars created, not lead
count.
    Identify: top 3 pipeline-creating channels, top content assets in winning
journeys,
    underperforming spend (high impression/click cost, low pipeline contribution).
    Then write a Slack-ready weekly brief in this format:
    - Pipeline Created This Week: $X
    - Top Channels (by $ sourced): 1. X ($Y) 2. X ($Y) 3. X ($Y)
    - Winning Content: [asset] appeared in [N]% of closed-won journeys
    - Underperforming Spend: [channel/campaign] generating leads not converting to
pipe
    - One insight or anomaly worth flagging to the revenue team
    Data: {{JSON.stringify($json.journeys)}}
  }]
}
```

5B. Parse and Post to Slack

Extract Claude's response text and post it to the designated Slack channel using an HTTP Request node with your Slack Bot token.

```
POST https://slack.com/api/chat.postMessage
Authorization: Bearer {{$credentials.slackBotToken}}
```

```
Content-Type: application/json

{
  "channel": "{{$env.SLACK_REVENUE_CHANNEL_ID}}",
  "text": "📊 *Marketing Pipeline Brief — Week of {{$now.toFormat('LLL d')}}*",
  "blocks": [{
    "type": "section",
    "text": { "type": "mrkdwn", "text": "{{claudeResponseText}}" }
  }]
}
```

5C. Auto-Tag Deals in CRM

As a final step, write Claude's channel attribution back to each deal in CRM as a custom property. This creates a queryable attribution field your team can use in reports.

```
PATCH https://api.[YOUR-CRM-API-URL]/.com/crm/v3/objects/deals/{{dealId}}
Authorization: Bearer {{$credentials.CRMApiKey}}

{
  "properties": {
    "marketing_source_attributed": "{{topChannelForDeal}}",
    "pipeline_journey_score": "{{claudeConfidenceScore}}"
  }
}
```

The Complete Workflow Architecture

Marketing-Sourced Pipeline — Full Node Chain

```
TRIGGER
└ Schedule: Every Monday 6:00 AM

DATA COLLECTION
└ CRM Deals Search (7-day window)
  └ Loop: For each deal
    └ CRM Contact Timeline (Engagements API)
└ Fibbler Account Engagement (7-day window)
  └ CRM Company Match (domain lookup)
└ Riverside Session Attendees
  └ CRM Contact Match (email lookup)

PROCESSING
└ Function: Stitch + sort journey objects by contact
└ Function: Merge Fibbler signals to matched contacts
└ Function: Append Riverside engagement events
└ Function: Build deal-level journey summary

AI ANALYSIS
└ Claude API: Score channels + generate weekly brief

OUTPUT
└ Slack: Post brief to #revenue-pipeline channel
└ CRM: PATCH deal records with attribution tags
```

Node Count and Estimated Build Time

Component	Nodes / Steps	Estimated Build Time
CRM data pull	4 nodes (trigger, search, loop, contacts)	1.5 hours
Fibbler integration	3 nodes (HTTP, match, merge)	1 hour
Riverside integration	3 nodes (HTTP, match, timeline write)	1 hour
Journey stitching logic	4 function nodes	2 hours
Claude prompt + parsing	2 nodes (HTTP, function)	1 hour
Slack + CRM output	2 nodes (Slack, PATCH)	30 minutes
Testing + QA	—	2–3 hours
Total	18 nodes	9–11 hours

Human-in-the-Loop Review

This agent automates data collection and analysis, but the revenue team should review Claude's output before acting on it. Run through this checklist each week when the Slack brief arrives.

1. **Verify deal count matches CRM.** Open CRM and confirm the number of deals created this week matches what Claude reported. Flag discrepancies to the ops team.
2. **Sanity-check the top channel.** If Claude names a channel you weren't running this week, there may be a UTM mismatch or attribution error. Investigate before sharing the brief upward.
3. **Review underperforming spend.** Claude flags campaigns with high lead volume but low pipeline contribution. Verify this with your paid media team before making budget changes.
4. **Check Fibbler match rate.** Low match rates (under 40%) indicate domain data hygiene issues in CRM. Run a company record audit if the rate drops two weeks in a row.
5. **Confirm Riverside attendance data freshness.** If no webinar ran this week, verify the Riverside pull returned zero records gracefully rather than erroring out.
6. **Review the flagged insight.** Claude includes one anomaly or pattern worth noting. Share this in your weekly revenue meeting with context.

Troubleshooting Common Issues

Problem	Likely Cause	Solution
Deals missing from brief	CRM search filter too narrow	Check createdate filter timezone offset; CRM timestamps are UTC
Fibbler returns no matches	Company domains not in CRM	Run a company enrichment pass; confirm domain field is populated
Riverside pull errors	Session ID not found	Confirm the sessionId is from a completed session, not a live one
Claude returns generic output	Journey data truncated	Reduce deals per run to 50; summarize journey before sending to Claude
Slack message not posting	Scopes missing on Bot token	Add chat:write and chat:write.public to your Slack app's OAuth scopes
CRM PATCH fails	Custom property not created	Create the marketing_source_attributed property in CRM Settings first
Attribution shows '(not set)'	UTMs not passing to CRM forms	Audit landing page form embed code; add hidden UTM fields

What’s Next

Agent 4 is the marketing layer of your revenue system. On its own, it tells you which channels are generating real pipeline. But its full power comes when it feeds into Agent 3 (Pipeline Visibility), which uses that channel data to classify whether deals are healthy or stalled — and which then feeds Agent 1 (Revenue Forecast) with a pipeline number you can actually make hiring and spend decisions on.

Once all four agents are running together, you have a live, self-calibrating revenue model. No gut feel. No stale spreadsheets. Just a number the whole team trusts.

Up next: Connect Agent 4 output to Agent 3 (Pipeline Visibility) to classify your live pipeline by health status using marketing signal data.